

Where Do Events Come From?

Reactive and Energy-Efficient Programming From The Ground Up

Anonymous Author(s)

Abstract

In reactive and event-based systems, execution is guided by an external environment that generates inputs to the application and consumes outputs from it. Reactive languages provide dedicated syntax and semantics to deal with events and greatly simplify the programming experience in this domain. Nevertheless, the environment is typically prefabricated in a host language and the very central concept of events is implemented externally to the reactive language. In this work, we propose an interrupt handler primitive for a reactive language targeting embedded systems in order to take control of the whole event loop, from input to output. We propose the new asynchronous primitive in the context of the synchronous language CÉU and discuss how they synergize to avoid race conditions at compile time, support lexically-scoped drivers, and provide automatic standby for applications.

Keywords interrupt service routines, reactive programming, standby, synchronous/asynchronous execution

1 Introduction

Reactive applications interact continuously and in real time with the external world through hardware peripherals such as sensors and actuators (e.g., buttons, displays, timers, etc.). These interactions are typically represented in software as input events flowing from the peripherals to the application and as output events flowing from the application to the peripherals. Peripherals can be encapsulated as a single component, *the environment*, which connects with the application through an event loop: the application sits idle until the environment awakes it on the occurrence of an input; the application reacts to the input and generates back outputs which affect the environment; the application becomes idle and the loop restarts.

Reactive languages provide dedicated syntax and semantics to deal with events and greatly simplify the development of applications in this domain. However, the environment is typically implemented in a host language (e.g., C) and controls the main event loop, invoking entry points into the reactive language runtime on the occurrence of inputs, and also receiving output calls from it. As examples, Esterel [2] relies on C for passing

events between the environment and the running program [9], while Elm [4] uses the concept of *ports*, which allows sending out values to JavaScript as commands and listening for values as subscriptions [3].

The event-based interface between the reactive language and the environment is arguably inevitable and also happens at the appropriate layer, since it connects the application with the operating system (OS) resources through a non-invasive API. However, this separation reveals the environment as a rigid system component that evolves in separate from the application. It also requires two languages and the programmer has to deal with multiple syntaxes, incompatible type systems, and different address spaces. Furthermore, in the context of embedded systems, a proper host OS may even be absent or lacking enough device drivers, which requires more low-level intervention from the application.

In this work, we propose interrupt service routines (ISRs) as an asynchronous primitive for the synchronous reactive language CÉU [10]. ISRs empower reactive applications to also implement their own device drivers and self-generate inputs, bypassing the need for a host environment. CÉU targets resource-constrained architectures, such as Arduino-compatible microcontrollers, in which applications run in the bare metal, without OS support. Applications can share data buffers with drivers to avoid unnecessary copying, with some static guarantee that no data races will occur. CÉU provides a lexical finalization mechanism that we adopt in drivers to properly disable interrupts and turn off peripherals completely. Finally, the synchronous semantics of CÉU enforces that applications react to inputs in bounded time and remain in idle states susceptible to standby. With the help of drivers and a power manager, we can put the microcontroller to sleep automatically at optimal sleeping modes after each input reaction. We implemented drivers for a variety of peripherals, such as GPIO, A/D converter, USART, SPI, and nRF24L01 transceiver. The applications built on top of these drivers show optimal energy savings due to automatic standby.

Our work is largely inspired by TinyOS [6], a low-power OS for wireless devices, which also provides asynchronous events with compile-time data race detection [5], and automatic energy management in idle CPU states [7]. We adapted these ideas to the structured reactive programming model of CÉU, and refined them with stronger

guarantees due to its support for lexically scoped lines of execution with automatic finalization.

2 Céu: Structured Synchronous Reactive Programming

CÉU [10] is a Esterel-based reactive programming language targeting resource-constrained embedded systems, with the characteristics that follow:

Reactive: code only executes in reactions to events and is idle most of the time.

Structured: programs use lexically-scoped structured control mechanisms such as `spawn` and `await` (to create and suspend lines of execution).

Synchronous: reactions run atomically and to completion on each line of execution, i.e., there's no implicit preemption or real parallelism.

Structured reactive programming lets developers write code in direct style, recovering from the inversion of control imposed by event-driven execution [8].

```

1 output high/low LED;
2 input high/low BUTTON;
3 input Packet RADIO_RECV;
4 par/or do
5   var high/low v = await BUTTON until (v==low);
6 with
7   finalize with
8     emit LED(low);
9   end
10 loop do
11   emit LED(high);
12   await RADIO_RECV;
13   emit LED(low);
14   await RADIO_RECV;
15 end
16 end

```

Listing 1. A Céu program that toggles the state of the LED whenever a radio packet is received, terminating on a button press, always with the LED off.

Listing 1 illustrates the main characteristics of Céu, namely event-driven I/O, lexically-scoped compositions, and synchronous execution. The program toggles the state of the LED whenever a radio packet is received, terminating on a button press, always with the LED off. The program first declares the LED output, and the BUTTON and RADIO_RECV input events (ln 1–3). The declarations include a payload type, i.e., each event occurrence carries a value of that type (`high/low` is a boolean type). Then, the program uses a `par/or` composition (ln 4–16) to run two lines of execution, aka *trails*, in parallel: a single-statement trail that waits for a button press before

terminating (ln 5), and an endless loop that toggles the LED on and off whenever a radio packet is received (ln 10–15). The `finalize` clause (ln 7–9) ensures that, no matter how its enclosing trail terminates (e.g., from a button press), the LED will be unconditionally turned off (ln 8).

All communication between the application and the environment is done through the `await` and `emit` primitives, which awaits an input event and generates an output event, respectively.

The `par/or`, which stands for *parallel-or*, is a lexical composition that terminates as soon as one of the trails terminates, but which also automatically aborts and finalizes the other trail(s). In the example, when the button is pressed (ln 5), not only the toggling loop will be aborted (ln 10–15), but the `finalize` clause will turn the LED off unconditionally (ln 8), since its enclosing block goes out of scope. The `par/or` is regarded as an *orthogonal preemption primitive* [1] because the trails need not to be tweaked in order to affect each other.

The synchronous execution model of Céu dictates that reactions to input events are atomic and that incoming events are never lost. The environment can only generate a new input after the application yields control and closes the event loop. In Listing 1, even if two packets arrive simultaneously in the environment, the synchronous model ensures atomicity and responsiveness in the application: the first `await` awakes and turns the LED off atomically (ln 12–13); and the second `await` will awake in sequence (ln 14).

2.1 Lexically-Scoped Interrupt Service Routines

Interrupts service routines (ISRs) are software entry points that execute in response to hardware interrupts from peripherals such as timers and GPIOs. ISRs are at the lowest interface layer between the hardware and software and are the absolute source of inputs to programs. An ISR starts to execute as soon as the hardware interrupt occurs, suspending the ongoing program flow abruptly. Such asynchronous behavior reflects the inherent concurrent nature of peripherals interacting with the external world.

We propose to extend Céu with asynchronous ISRs. However, asynchronous execution confronts the synchronous mindset of Céu since the assumption that reactions are atomic no longer holds. Not only this may lead to race conditions at a fine grain, but may also affect the ordering of events at a coarse grain: the effect of an earlier event may be perceived after the effect of a later event. Still, our goal is to preserve the well-behaved interaction between the Céu application and the environment even in the presence of asynchronous ISRs. Our idea is to push all subtleties of asynchronous execution into device

```

1 output high/low PIN_13; // connected to an LED
2 input high/low PIN_02; // connected to a button
3 #include "gpio.ceu" // GPIO driver
4
5 emit PIN_13(low);
6 loop do
7   var high/low v = await PIN_02;
8   emit PIN_13(v);
9 end

```

Listing 2. Synchronous application that turns the LED on whenever the button is pressed and off whenever it is unpressed.

```

1 // OUTPUT DRIVER
2
3 { pinMode(13, OUTPUT); }
4 output (high/low v) PIN_13 do
5   { digitalWrite(13, @v); }
6 end
7
8 // INPUT DRIVER
9
10 input high/low PIN_02;
11 {
12   pinMode(2, INPUT_PULLUP);
13   EICRA |= (1 << ISC00); // sets INTO
14   EIMSK |= (1 << INTO); // turns on INTO
15 }
16 spawn async/isr [INT0_vect] do
17   emit PIN_02({digitalRead(2)});
18 end

```

Listing 3. Asynchronous driver for GPIO which, on hardware interrupts, emits an input event into a queue.

drivers, which emit input events to regular synchronous code in CÉU, exactly as before.

As a first example, Listing 2 and 3 uses GPIOs to connect a button (pin 02) to an LED (pin 13) via software such that the LED is on whenever the button is pressed and off whenever it is unpressed.

Listing 2, the synchronous side, first declares its interface with the external world and includes the asynchronous side as a driver (ln 1–3). It then starts with the LED off (ln 5) and enters a loop (ln 6–9) that, whenever the button changes (ln 7), toggles the state of the LED with the new value (ln 8).

Listing 3, the asynchronous side, implements the driver for the output and input events. An `output` implementation (ln 4–6) is similar to a parameterized subroutine: whenever the application invokes `emit`, the output body executes atomically. CÉU supports inline C between curly braces (ln 3,5) with interpolation to evaluate CÉU

expressions (e.g., `@v`). This allows drivers to take advantage of existing libraries of embedded toolchains, such as Arduino. In the example, when the driver is included, it configures pin 13 as output (ln 3) and sets its new state whenever PIN_13 is emitted (ln 5). An `input` event implementation uses an ISR registered with the `spawn async/isr` primitive (ln 16–18), which is automatically invoked whenever the associated interrupt occurs (e.g., `INT0_vect`). An ISR in CÉU will typically emit an input event inside its body to awake the synchronous side (ln 17). However, although the ISR executes asynchronously when the interrupt occurs, the `emit` goes into a queue (to be discussed) and does not interrupt the ongoing (suspended) reaction on the synchronous side. In the example, the input driver also configures pin 2 to behave as input and to generate external interrupts on level transitions (ln 11–15).

The example illustrates the clear separation between the application and the driver through an event-driven interface. Now, the whole application is written in CÉU: the synchronous side remains well behaved with no low-level calls, and the asynchronous side deals with the complexity of device drivers. Nevertheless, drivers are typically write-once components developed by embedded specialists, which are reused in most applications. Note that each supported architecture requires a mapping between the `async/isr` and the actual interrupt vector table (which we provide for the *AVR/ATmega* and *ARM/Cortex* microcontrollers).

The example in Listing 4 illustrates the use of a lexically-scoped driver for the A/D converter. Note that the driver (ln 1–14) and application (ln 16–27) are actually in separate files. The application is a typical periodic sensor sampling loop, but which designates an explicit `do-end` block (ln 20–24) to include the driver (ln 21) and use it (ln 22–23). The driver specifies a finalization code (ln 6–12) which executes unconditionally whenever its enclosing block goes out of scope (ln 24). The finalize completely disables all A/D functionality to save energy. As discussed in the previous section, the finalizer would also execute if aborted from a `par/or` enclosing the driver. Note that the driver interface (ln 3–4) is visible only inside the block and thus cannot be inadvertently used outside it.

2.2 Safe Shared-Memory Concurrency

In CÉU, when multiple trails awake in the same reaction, they execute in lexical order, i.e., in the order they appear in the source code. In Figure 1, both examples define a shared variable (ln 3) and assign to it in parallel trails (ln 6, 9). In Listing 5, the two assignments to `x` can only execute in reactions to different events **A** and **B** (ln 5,8), which cannot occur simultaneously in accordance with the synchronous model. In Listing 6, the two assignments

```

331 1 // A/D DRIVER (adc.ceu)
332 2
333 3 output int ADC_REQUEST;
334 4 input int ADC_DONE;
335 5
336 6 do finalize with
337 7 {
338 8     ADCSRA &= B01111111; // disables A/D converter
339 9     ACSR = B10000000; // disables comparator
340 10    DIDRO |= B00111111; // disables pins
341 11 }
342 12 end
343 13 ... // driver initialization
344 14 ... // input / output implementations
345 15
346 16 // APPLICATION (main.ceu)
347 17
348 18 loop do
349 19     var int v;
350 20     do
351 21         #include "adc.ceu" // driver contents above
352 22         emit ADC_REQUEST(A0);
353 23         v = await ADC_DONE;
354 24     end
355 25     ... // uses sensor value "v"
356 26     await 1h;
357 27 end

```

Listing 4. The A/D driver (ln 1–14) is included in the application inside a lexically-scoped block (ln 20–24), which turns off all driver functionalities automatically on termination.

<pre> 362 input void A; 363 input void B; 364 var int x = 1; 365 par/and do 366 await A; 367 x = x + 1; 368 with 369 await B; 370 x = x * 2; 371 end </pre>	<pre> 1 2 3 4 5 6 7 8 9 10 </pre>	<pre> 362 input void A; 363 // (empty line) 364 var int y = 1; 365 par/and do 366 await A; 367 y = y + 1; 368 with 369 await A; 370 y = y * 2; 371 end </pre>
---	-----------------------------------	---

Listing 5. Shared x **Listing 6.** Shared y

Figure 1. Accesses to shared x never concurrent. Accesses to shared y are concurrent but still deterministic.

to y are simultaneous because they execute in reaction to the same event A. Nevertheless, because CÉU follows lexical order, the execution is still deterministic, and y always becomes 4 ((1+1)*2). Even so, CÉU performs (optional) concurrency checks at compile time to detect conflicting accesses to shared variables: if a variable

```

386 1 // USART INTERFACE
387 2
388 3 input none USART_RX;
389 4 var[32*] byte rx_buf; // '*' = circular
390 5
391 6 // DRIVER
392 7
393 8 spawn async/isr [USART_RX_vect] do
394 9     rx_buf = rx_buf .. [{UDR0}];
395 10    emit USART_RX;
396 11 end
397 12
398 13 // APPLICATION
399 14
400 15 loop do
401 16     await USART_RX;
402 17     var int i;
403 18     loop i in [0 -> $rx_buf[ do // '$' = length ($)
404 19         // uses rx_buf[i]
405 20     end
406 21     $rx_buf = 0;
407 22 end

```

Listing 7. The USART buffer is shared between the ISR and the application, resulting in a potential race condition. (The driver and application code are in separate files.)

is written in a trail segment, then a concurrent trail segment cannot access that variable [10].

The addition of asynchronous ISRs poses real threats concerning race conditions since they interrupt programs at arbitrary points. Listing 7 illustrates a minimum USART application (ln 13–22) that consumes incoming bytes (ln 18–21) as they arrive (ln 16). The USART API (ln 1–4) exposes an input event to signal incoming data (ln 3) and a circular buffer to prevent data loss (ln 4). The buffer is indispensable because the microcontroller may cope with the USART speed. The driver ISR (ln 6–11) simply appends incoming bytes to the end of the buffer (ln 9) and signals the application (ln 10). As a possible race condition, note that the ISR may fire and append a new byte to the buffer (ln 9) just before the application clears it (ln 21), in which case the new byte will be lost forever.

CÉU treats an `async/isr` as a block of code that runs in parallel with the rest of the program (hence the required prefix `spawn`). This way, it is clear for the compiler that the accesses to `rx_buf` in Listing 7 may lead to a race condition (ln 9,21), and CÉU raises a compile-time error. The programmer is responsible to protect concurrent memory accesses with `atomic` blocks, which disables interrupts (supposedly) for a short period of time.

However, we do not expect that application programmers should be required to resolve race conditions. CÉU

```

441 1 // USART INTERFACE
442 2
443 3 code Usart_RX (var&[] byte buf, var int n) -> none;
444 4
445 5 // DRIVER
446 6
447 7 input none USART_RX;
448 8 var[32*] byte rx_buf;
449 9
450 10 spawn async/isr [USART_RX_vect] do
451 11     rx_buf = rx_buf .. [{UDRO}];
452 12     emit USART_RX;
453 13 end
454 14
455 15 code Usart_RX (var&[] byte buf, var int n) -> none
456 16 do
457 17     loop do
458 18         atomic do
459 19             buf = buf..rx_buf;
460 20             $rx_buf = 0;
461 21         end
462 22         if $buf >= n then
463 23             break;
464 24         end
465 25         await USART_RX;
466 26     end
467 27 end
468 28
469 29 // APPLICATION
470 30
471 31 loop do
472 32     var[255] byte buf;
473 33     await Usart_RX(&buf, 10);
474 34     // uses buf
475 35 end

```

Listing 8. The USART driver now exposes a safer (and more friendly) interface to the application. *(The driver and application code are in separate files.)*

supports reactive abstractions (similar to coroutines) that help hiding the concurrency complexity inside the driver. Listing 8 changes the USART API (ln 1–3) to expose a code abstraction that expects a reference to a buffer and a number of bytes to receive. Now, the application (ln 29–35) invokes the abstraction (ln 33) passing a local buffer (ln 32). The driver (ln 5–27) now hides the low level interface of the previous example (ln 7–8) and implements the code abstraction (ln 15–27) that protects the access to the shared buffer with an `atomic` block (ln 18–21). The abstraction copies the driver buffer into the application buffer (ln 19) up to the requested number of bytes (ln 22). On the one hand the extra copying incurs a memory and runtime overhead, but on the other hand it prevents race conditions with the ISR.

```

496 1 evt_t queue[MAX];           // input queue
497 2 void main () {
498 3     ceu_start();             // "boot reaction"
499 4     while (1) {
500 5         evt_t evt;
501 6         if (ceu_input(&evt)) { // query input queue
502 7             ceu_sync(&evt);    // execute synchronous
503 8         } else {
504 9             ceu_pm_sleep();    // nothing to execute
505 10        }
506 11    }
507 12 }

```

Listing 9. Overall runtime architecture of CÉU with an input queue (ln 1), event loop (ln 4–11), synchronous execution (ln 7), and standby mode (ln 9).

Support for ISRs in the same language and memory space of the application allows programmers to choose between efficiency and safety during development by providing multiple interfaces with different levels of abstraction.

2.3 Energy-Efficient Runtime

Our runtime alternates between synchronous, well-behaved execution, and asynchronous, unpredictable ISRs. Since the language now has full control over the event loop and only executes from interrupts, it is possible for the runtime to enter in sleep mode automatically to save energy.

Listing 9 is a realization of this architecture: The runtime defines an input queue (ln 1) in which ISRs enqueue new events (e.g., Listing 8, ln 12). In the `main` function, we first generate the *boot reaction* (ln 3), which starts the program, spawns the ISRs, and reaches the first `await` statements in the multiple lines of code. Then, the runtime enters an infinite loop (ln 4–11) that queries the input queue (ln 6) to awake the program (ln 7). Note that ISRs execute asynchronously with the loop, but they typically only emit an input event into the queue which will be queried in a subsequent iteration. Note also that if the queue is empty, the runtime enters in sleep mode to save energy (ln 9), and will only awake on a new hardware interrupt.

Microcontrollers typically support multiple levels of sleep mode, each progressively saves more energy at the expense of keeping less functionality active. As an example, the least efficient mode of the *ATmega328P* allows for timer interrupts since it keeps its internal clock active, while the most efficient mode turns off all peripherals and can only awake from external interrupts.

Our runtime supports three compile-time configurations for `ceu_pm_sleep` (ln 9): The first configuration is a *nop* which simply keeps the event loop running all the

```

551 1 // POWER MANAGER (in C)
552 2
553 3 enum {
554 4     CEU_PM_ADC,
555 5     CEU_PM_TIMER1,
556 6     CEU_PM_USART,
557 7     ...
558 8 };
559 9
559 10 void ceu_pm_sleep (void) {
560 11     if (ceu_pm_get(CEU_PM_USART) || ...) {
561 12         sleep_mode_1(...);
562 13     } else if (ceu_pm_get(CEU_PM_ADC)) {
563 14         sleep_mode_2(...);
564 15     } else {
565 16         sleep_mode_3(...);
566 17     }
567 18 }
568 19
568 20 // USART DRIVER (in Ceu)
569 21
570 22 code Usart_RX (var&[] byte buf, var int n) -> none
571 23 do
572 24     {ceu_pm_set(CEU_PM_USART, 1);}
573 25 do finalize with
574 26     {ceu_pm_set(CEU_PM_USART, 0);}
575 27 end
576 28 ... // same as in Listing 7
577 29 end
578 30
578 31 // APPLICATION (in Ceu)
579 32
580 33 input high/low PIN_02; // connected to a button
581 34 loop do
582 35     par/or do
583 36         await PIN_02;
584 37     with
585 38         var[255] byte buf;
586 39         await Usart_RX(&buf, 10);
587 40         // uses buf
588 41     end
589 42     await PIN_02;
590 43 end

```

Listing 10. Interaction between the power manager, USART driver, and application (each code is actually in a separate file).

time without ever sleeping. This configuration is useful when introducing new platforms. The second configuration chooses a (inefficient) sleep mode that keeps all peripherals active. Although this configuration is not the most energy efficient, at least, it requires no extra assistance from the drivers. The third configuration keeps a bit vector of the active drivers during runtime and chooses the most efficient mode, querying this vector every time `ceu_pm_sleep` is called.

The third configuration achieves optimal efficiency but requires a tight interaction between the drivers and the power management runtime. Listing 10 unveils the power manager (ln 1–18), the modified USART driver (ln 20–29), and an illustrative application (ln 31–43). The application is a loop that awaits in parallel for a button press (ln 36) or receiving 10 bytes (ln 39). Let's call this initial state **STATE-A**. The **par/or** terminates when either of them occurs, going to the next line that awaits another button click (ln 42). Let's call this other state **STATE-B**. After the button click, the program loops back to **STATE-A**. While in **STATE-A**, the program can awake from USART and external interrupts, which means that the power manager should choose a sleep mode in which the USART remains operational. While in **STATE-B**, only external interrupts should awake the program, which means that the power manager may use the most efficient sleep mode. The power manager enumerates all microcontroller's peripherals (ln 3–8) and, before sleeping, queries their states (ln 11,13) to choose the most appropriate sleep mode (ln 12,14,16). The USART driver needs to be extended with calls to enable and disable the USART state inside the power manager: just before awaiting, the driver enables the USART (ln 24) and creates a finalization block to unregister it on termination (ln 25–27). Termination may occur either directly after receiving the requested number of bytes, or indirectly if the **par/or** in the application terminates on a button click (ln 36).

Note that applications need not to be explicit about energy management to take advantage of sleep modes. All happens automatically because of the synchronous-reactive execution model and energy-aware runtime of CÉU.

3 Conclusion

We extend the synchronous language CÉU with asynchronous interrupt service routines (ISRs) to take control of the whole event loop in reactive systems, from input to output. ISRs empower reactive applications to also implement their own device drivers and self-generate inputs, bypassing the need for a host environment. However, asynchronous execution confronts the well-behaved semantics of synchronous languages with possible race conditions. In order to mitigate this threat, we extend the static concurrency checks of CÉU to detect data races between ISRs and the application.

Our approach for writing drivers take advantage of the lexical structured mechanisms of CÉU, such as parallel compositions and abortion of lines of execution, to offer stronger guarantees. For instance, the visibility of a driver can be delimited to a scoped block with the guarantee that its resources will not leak after usage. The same mechanism permits that drivers signal the

power manager that its peripheral is no longer required, which allows applications to use optimal sleep modes to save energy automatically. Our initial tests show optimal energy savings for applications in idle states.

References

- [1] Gérard Berry. 1993. Preemption in Concurrent Systems.. In *FSTTCS (LNCS)*, Vol. 761. Springer, 72–93.
- [2] Frédéric Boussinot and Robert De Simone. 1991. The Esterel language. *Proc. IEEE* 79, 9 (Sep 1991), 1293–1304.
- [3] Evan Czaplicki. 2018. Elm and JavaScript Interop. <https://guide.elm-lang.org/interop/javascript.html> (accessed in Jul-2018).
- [4] Evan Czaplicki and Stephen Chong. 2013. Asynchronous functional reactive programming for GUIs.. In *Proceedings of PLDI'13*. 411–422.
- [5] David Gay et al. 2003. The nesC Language: A Holistic Approach to Networked Embedded Systems. In *Proceedings of PLDI'03*. 1–11.
- [6] Jason Hill et al. 2000. System architecture directions for networked sensors. *SIGPLAN Notices* 35 (November 2000), 93–104. Issue 11.
- [7] Kevin Klues et al. 2007. Integrating Concurrency Control and Energy Management in Device Drivers. In *Proceedings of SOSPP'07*. ACM, New York, NY, USA, 251–264.
- [8] Ingo Maier, Tiark Rompf, and Martin Odersky. 2010. *Deprecating the observer pattern*. Technical Report.
- [9] Dumitru Potop-Butucaru, Stephen A Edwards, and Gerard Berry. 2007. *Compiling esterel*. Vol. 86. Springer Science & Business Media.
- [10] Francisco Sant'Anna et al. 2013. Safe System-level Concurrency on Resource-Constrained Nodes. In *Proceedings of SenSys'13*. ACM.